case-study-brief

Product consultant case study

Rakamin, capability-based workforce transformation

Role: Product Consultant
Submission window: 3 to 5 business days from receipt
Deliverables: a deck and a working prototype, both defined in Section 4
Live session: you present the deck to a director panel for 60 minutes, scheduled after submission

> This is a big case, drawn from work we actually do. We made it demanding on purpose, because the role is.
> Use AI tools freely throughout. There is no single correct answer. We read your thinking and your judgment, and whether this is work you would be comfortable presenting to a director.

Section 1. Who we are

Rakamin is a workforce architecture and AI-native consulting firm working with large Indonesian enterprises: state-owned enterprises, banks, and national corporates. We run a hybrid consulting and product model. We go deep into a client's workforce and HR data, redesign the architecture underneath it, build AI on top, and hand their leaders decision intelligence they can act on.

Our 2026 thesis rests on one shift:

> The era of the fixed job title is ending. AI is changing what a role *is* faster than any org chart can be redrawn. Work that took three weeks now takes two hours. Enterprises that bind their people to static positions cannot move at that speed.

So we help clients move from a position-based workforce to a capability-based one. In a position-based workforce, a person is a title and a seat on the org chart. In a capability-based one, a person is a grade, a domain, and a portfolio of skills that can be redeployed as the work changes. We built this architecture on ourselves first: job families, leveling, capability bands, skill taxonomy, project economics. Now we install it for clients.

A Product Consultant leads that engagement in front of the client's directors. That is the role you are interviewing for.

What the role is

When a consultant leads a digital transformation, the work does not split into "talking to the client" and "doing the technical part." It is one job. You move across four pillars and they have to hold together:

Pillar	What you do
Data	Make sense of a workforce data mess across systems that do not agree with each other. Resolve identities, normalize taxonomy, quantify the workforce, model the economics.
Design	Design the capability architecture itself and the systems and AI that run it. Reshape the business process, not only the org chart.
Decision	Turn a vague board mandate into a sequenced bet. Cut scope honestly. Make a recommendation you can defend with numbers in front of a director.
Digital transformation	Move thousands of people from positions to capabilities without renegotiating thousands of contracts or triggering industrial action. Lead the adoption, the resistance, and the politics.

Your two deliverables exercise all four at once.

Section 2. The scenario

The client

The client is masked. Everything below is a composite of a real engagement. We call the work Project Meridian and the client the Client.

The Client is a state-owned energy enterprise:

- ~52,000 employees, across 6 regional operating units and a head office
- Regulated, politically visible, reports to a Ministry and a Board of Commissioners
- A 60-year-old institution with a deep seniority culture. Promotion has meant waiting for a seat to open ("menunggu kursi kosong"), not demonstrating new capability

Their workforce architecture is position-based:

- 1,800+ unique job titles, each bound to a grade that is bound to a box on the org chart
- One generic 5-level competency model applied to every role, from field technician to division head
- A person's identity, pay, and authority all flow from their title. Change the title and you have, in practice, reopened their employment contract

Their HR and workforce technology stack:

System	Vendor	Known issue
HRIS	SAP SuccessFactors	Incomplete historical data, low manager adoption
Competency and assessment	Legacy in-house system	Generic 5-level model, not role-specific, rarely updated
LMS	Moodle	Poor skill taxonomy, inconsistent course categories
Performance	Custom per-region tools	Rating scales differ across the 6 regional units
Payroll	Oracle	Restricted access under Ministry data-governance policy

More data reality to assume:

- 45% of employee records have missing or blank skill fields
- Job titles are inconsistent across regions. An "Engineer Madya" in one unit maps to a "Senior Technician II" in another
- Employee IDs differ across systems (HRIS: `EMP-48213`, LMS: `learner_9981`, Payroll: `PN-7740`)
- Performance data before 2023 is unreliable and was scored on scales that have since changed

The trigger

The Board has mandated a Human Capital Transformation for the AI era. The Director of Human Capital owns it, and she has a hard deadline: present a credible plan and the first proof points to the Board and the Ministry in 90 days.

She has told you three things directly:

> "I need to know which of our people the next wave of AI will displace, and what we do with them."
> "I cannot do mass PHK (layoffs). Politically and socially, it is not on the table. The Ministry will not allow it."
> "I am not replacing SAP or Oracle. Do not bring me a system migration. Bring me a way to run my workforce differently on top of what I already have."

Her open questions:

1. Which job families are most exposed to AI disruption in the next 2 to 3 years, and how exposed?
2. How do we move 52,000 people from positions to capabilities without renegotiating 52,000 contracts and without triggering the union?
3. How do we tell, with evidence, who can be reskilled and redeployed and who is surplus?
4. How do we prove the return on this to a Board and a Ministry that think in budgets, not in competency models?

The engagement: Rakamin has been contracted for Rp 3.5B over 12 months. You are the Product Consultant leading it. The first Board milestone is in 90 days.

Section 3. The complication (read carefully, this is the case)

On Day 30, your data work turns up something that changes the engagement.

> One job family, Field Metering and Manual Operations (~6,000 people, ~11% of the workforce), is roughly 70% automatable within 18 months, through a mix of smart-meter rollout, IoT, and AI-assisted operations the Client has already started procuring in a separate division. Word reaches the labor union (Serikat Pekerja) that "consultants are mapping who gets cut."

You now hold a contradiction:

- The data says ~4,200 roles' worth of work in that family is going away.
- The Director cannot present a layoff. She has to present a redeployment the Board believes, the Ministry approves, and the union does not strike over.
- A second constraint arrives at the same time. Legal confirms Payroll (Oracle) is off-limits under the Ministry's data-governance policy. You will not get compensation data for this engagement. Any capability mapping or workforce economics you build has to work without it.

Your design, your numbers, your change plan, and the Day 90 story all have to absorb this. A plan that skips it does not meet the brief.

Section 4. What you must deliver

Two deliverables, not two separate exercises. The deck makes the case. The prototype lets us experience the same solution. We read how well the two hold together.

Deliverable 1, the deck

The board-grade story you would put in front of the Director and, through her, the Board and the Ministry. You present it live (Section 5). Slide count is your call. Clarity beats volume.

The four the deck must cover

1. **The problem.** "AI transformation" is a mandate, not a problem. State the three or so focus areas you considered and why you chose this one over the others.
2. **The urgency.** Why now, in numbers. The AI-disruption clock on the metering family, the cost of a year of standing still, the 90-day Board deadline. Make the room feel the cost of waiting.
3. **The business outcome.** What success looks like in the Board's language: a workforce that can be redeployed as work changes, displacement converted to redeployment, the return on the Rp 3.5B. Outcomes the Board can be held to, not aspirations.
4. **The solution.** The capability architecture and the move from position-based to capability-based, in a form the Board can repeat back. Items 7 and 10 carry the detail.

What else it must carry

5. **The strategic bet and the scope cut.** What you sequence first, and what you are not solving in 90 days. If nothing is out of scope, you have not decided anything.
6. **Where your architecture comes from.** Capability-based workforce design is not new. Large enterprises have been installing it for years, some well, some expensively and badly, and established frameworks and reference models sit behind that work.

   Show us you know what already exists before you design anything. What the options are, who runs them, and what each one costs an organization of this size to adopt. Then take a position: adopt something proven as it comes, build a structure specific to this Client, or mix the two on purpose. All three are defensible, so name what yours costs you. A well-argued original design beats a catalogue of frameworks, but a design built without knowing the field is a guess.
7. **The capability architecture, shown.** Job families, leveling, capability bands, skill taxonomy, grounded in established workforce theory and named explicitly. Show how 1,800+ titles collapse into a defensible structure and how a person becomes grade plus domain plus skill portfolio. Show one business process in the affected family before and after. Workforce here means reshaping the work, not relabeling people.
8. **The chain of reasoning underneath it.** Every element of the architecture has to trace back to something, and we will follow the trace in both directions.

   One way. What makes this the right set of capabilities for this Client and this business, rather than a list that would fit any energy company? What decides that a capability sits at one level and not the next one up?

   The other way. What establishes that a specific person holds a given capability, at a given level? Remember what you are working with. 45% of skill records are blank, and the pre-2023 performance data cannot be trusted. Whatever mechanism you propose, be ready to say how far you trust it, how you would know when it is wrong, and what the Director is allowed to decide on a signal that weak.

   Boards kill transformations over a single number nobody can trace.
9. **The data insight.** What the workforce data reveals: the AI-exposure map by job family, reskillability, where the risk concentrates. Anyone can describe a workforce. Show us you can read one.
10. **The migration and change plan.** How you move 52,000 people in sequence, how you handle the "my title is my identity" problem in a seniority culture, and the union and Ministry engagement plan. When you bring them in, and with what.
11. **The complication, answered.** Turn the ~4,200 displaced-work roles into a redeployment story the Director can present and the union will not strike over. Name the adjacent capability bands the metering people move into, your evidence it is real, and what you do with the people for whom redeployment will not work.
12. **The 90-day roadmap.** Day 30, 60, 90 milestones, each one a demo-able output rather than a phase like "data collection." What does the Board see at each gate?
13. **The economics.** The numbers a Board can defend: cost of inaction against cost of transformation, redeployment against the layoff-then-rehire it cannot do anyway, and the engagement ROI. State your assumptions where you lack data. You will lack data, because payroll is off-limits. Show the working behind the numbers.
14. **Risks and repeatability.** What could go wrong and how you mitigate it. Then the question we always ask: if Project Meridian works, how does this become a reusable Rakamin capability-architecture method for the next five clients instead of a bespoke build for one?

Deliverable 2, the working prototype

A working desktop web application, built with AI, that runs and solves the problem. The interface has to be good enough to put in front of an executive. A static mockup is not enough.

What we read here is the experience, not a data pipeline. Give us the user journey and the UX.

What it must do

- **Be a single self-contained HTML file**, content baked in. It opens and runs in any browser in one click. Section 7 has the hosting mechanics.
- **Tell a user journey end to end.** Pick the primary user, the Director or an HR manager or an affected employee, and let us click through how they use the solution, from the problem they walk in with to the decision they walk out with.
- **Make the problem visibly solved.** By the end of the journey the user can see and act on something they could not before: the AI-exposure picture, the capability map, the redeployment path. The solution should be obvious in use, not explained in a caption.
- **Let us see what the screen is asserting.** When the interface puts a level, a score, or a recommendation in front of the user, we should be able to reach what it rests on and how sure it is. An executive who cannot tell a measured number from an inferred one will trust all of it or none of it. Both are failures of the interface.
- **Engineer the AHA moment.** One specific moment where the client sees their workforce in a way they never have and the value lands. Design for it on purpose, and be ready to say why it works.
- **Make it feel real.** Use realistic, workforce-shaped sample content. A few dozen believable records carry the journey further than a full dataset.

Design taste matters, and it is the thing we are most sensitive about. A Product Consultant directs designers and engineers. You do not have to be one, but a good eye and the judgment to guide it both count. We look for intentional hierarchy, calibrated information density, honest handling of confidence and uncertainty, and a consistent visual language.

We do not accept AI slop. The prototype has to represent your own standard for UI and UX, not whatever the model produced on the first pass. It is the only evidence of your taste we get, so we read an unedited generation as your bar. Build it with AI, then keep pushing until every screen is one you would defend in front of the Director.

Use AI coding tools (Claude, Cursor, and the like) to build this. It is not a test of coding ability, it is a test of whether you can turn judgment into a working solution with the tools available today.

Section 5. The live presentation (director panel)

After submission you present the deck in a 60-minute session. The panel plays three people: the Director of Human Capital, a Board member focused on budget and risk, and a union representative. The prototype is open on the table, and we will probably ask you to walk us through it.

Expect the panel to push. This is where we see how you:

- Hold the room and lead the conversation at director level
- Stand behind a number when it is questioned, and explain your reasoning
- Hear the union's concerns and answer them without retreating into jargon
- Change your position when a point lands, without losing the thread

Two things we will test specifically

The trade-offs, stated plainly. For every significant choice we will ask what the alternatives were, what each one would have cost you, and why you took this one anyway. That covers the focus area you picked, the architecture you adopted or built, the scope you cut, the sequence of the 90 days, the primary user of the prototype, the assumptions holding up your economics. Be ready to say what your recommendation is worse at, not only what it is good at. We will keep pushing until we find the cost you did not name, so naming it yourself first is the stronger move.

Full ownership of everything you submit, including what AI produced. We may point at any number, any slide, or any screen in the prototype and ask you to derive it in front of us. "That is what the model gave me" is not an answer. So go through the deliverables yourself before you submit. Numbers reconcile across slides. Every figure traces to a source or to an assumption you have stated. Every claim you make about workforce theory is one you can explain. The screens we click behave the way the deck says they do. If we find a flaw you have not already caught and flagged, we read it as work you did not validate, and that costs you more than the flaw itself.

Section 6. Ground rules

This brief is everything we will provide for this stage. We will not share more materials or answer clarifying questions before you submit. Treat anything not stated here as yours to assume, and proceed. How you navigate the ambiguity is part of what we hope to learn, and we are glad to discuss any of it at the live session.

Use AI tools throughout, and validate what comes back. Note which tools you used and how. What we are reading is your judgment about when to use AI and where it should not be left to decide on its own. Whatever you submit is yours to defend, line by line, in the live session (Section 5).

Cite your assumptions. You will not have all the data, by design. When you assume, say what you assumed and why. Unstated assumptions are a common reason engagements go off track.

Be direct and take a clear position. A well-argued point of view we can engage with is more useful than a balanced view that avoids committing.

Section 7. Submission

Submit a single PDF through the hiring portal.

The PDF contains:

1. **The deck.** The full story from Section 4. Include your economics working, the assumptions and the model behind the numbers, inline or as an appendix.
2. **The prototype link.** A page near the front of the deck, or the very first page, carrying the live URL to your working prototype, written out in full as plain text so it is readable and clickable. Add a line or two telling us which screen is the AHA moment and where to click to reach it.

Upload the HTML file somewhere we can reach from that link. Netlify Drop, GitHub Pages, Vercel, a public Google Drive or Dropbox link, any host that serves a file. What matters is that we click the link and the prototype opens and runs, with no login, no build step, and no setup on our side.

Two things to check before you submit:

- **Open the link yourself in a private or incognito window.** If it asks us to sign in or request access, we cannot review it.
- **Keep it live through the live session.** We will open it again during the panel.

Close the deck with one short note, 3 to 5 sentences:

> **What is the one decision in your submission you are least confident about, and what would you need to validate it?**

We schedule the live presentation after we review your submission.